


Raju Gandhi

---

# PRACTICAL DOCKER

---



# RAJU GANDHI

   @LOOSELYTYPED  
FOUNDER - DEFMACRO SOFTWARE

# (POLL - Single Choice)

## Embracing DevOps

- We do it all — Ci/Cd pipelines, Everything-as-Code, Centralized log/event/monitoring etc
- We have a few pieces in the works
- Just getting started


**WHY?**

**BUILD ONCE, RUN ANYWHERE**

# WHY?

- Local application development and testing
- Team (and OSS) collaboration
- Ci/Cd
- Higher density deployments
- Deterministic


**IMAGES**








Image




Image

Container



Image

Container



Image

Container

# IMAGES

- The docker artifact
- Opaque
- Shared using a registry like <http://hub.docker.com/>

# IMAGES

- The "template" for containers
  - Figure out what you need "fixed" a.k.a "compile" time settings
  - Figure out what you want the consumer to specify a.k.a "runtime"

# (POLL - Single Choice)

## Docker usage


- Everywhere we can! — dev/qa/prod/Ci-Cd pipelines
- Not all the way to production yet (only development or lower tiers)
- Just testing the waters
- Not there yet

**WHAT IS A CONTAINER?**





# WHAT IS A CONTAINER

- An instance of an image
- With a r/w layer on top
- Configured with resource limits (cpu/memory), network settings and volume mounts etc on "create"



docker run





make changes



Writeable layer

Base Image

Host OS (Kernel)



Image



Container

docker commit



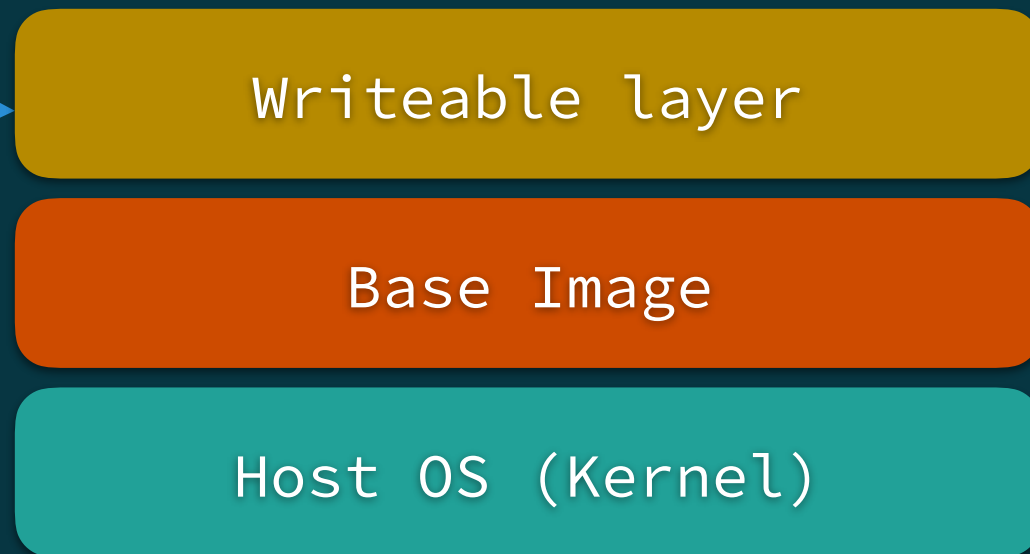
Writeable layer

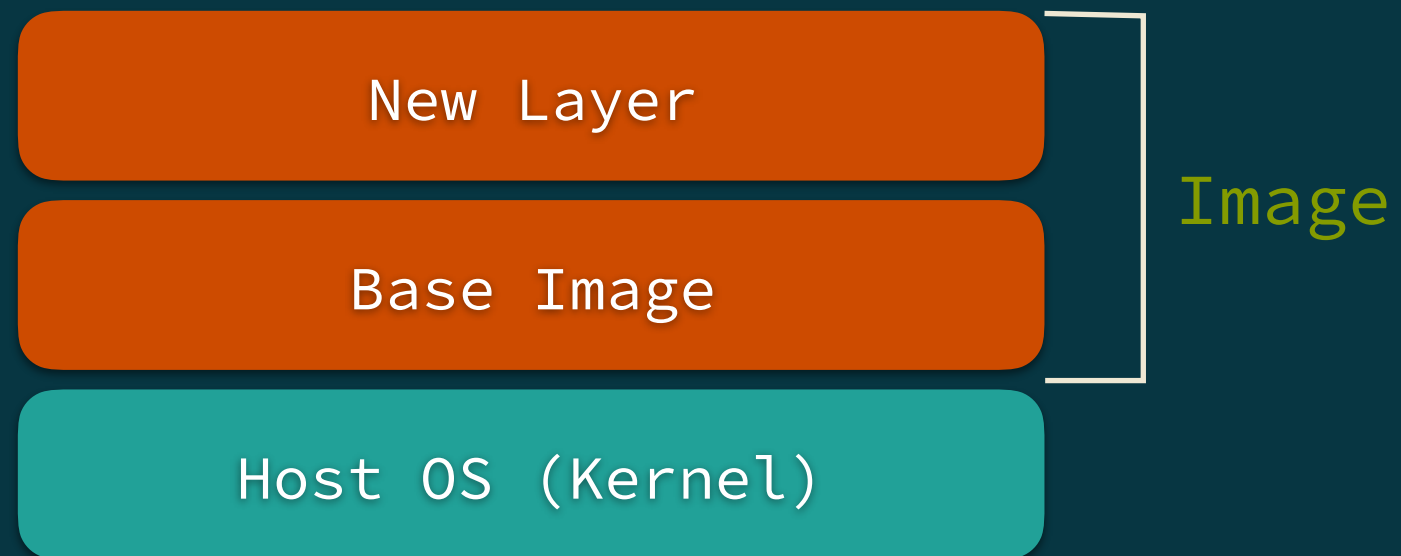
Base Image

Host OS (Kernel)

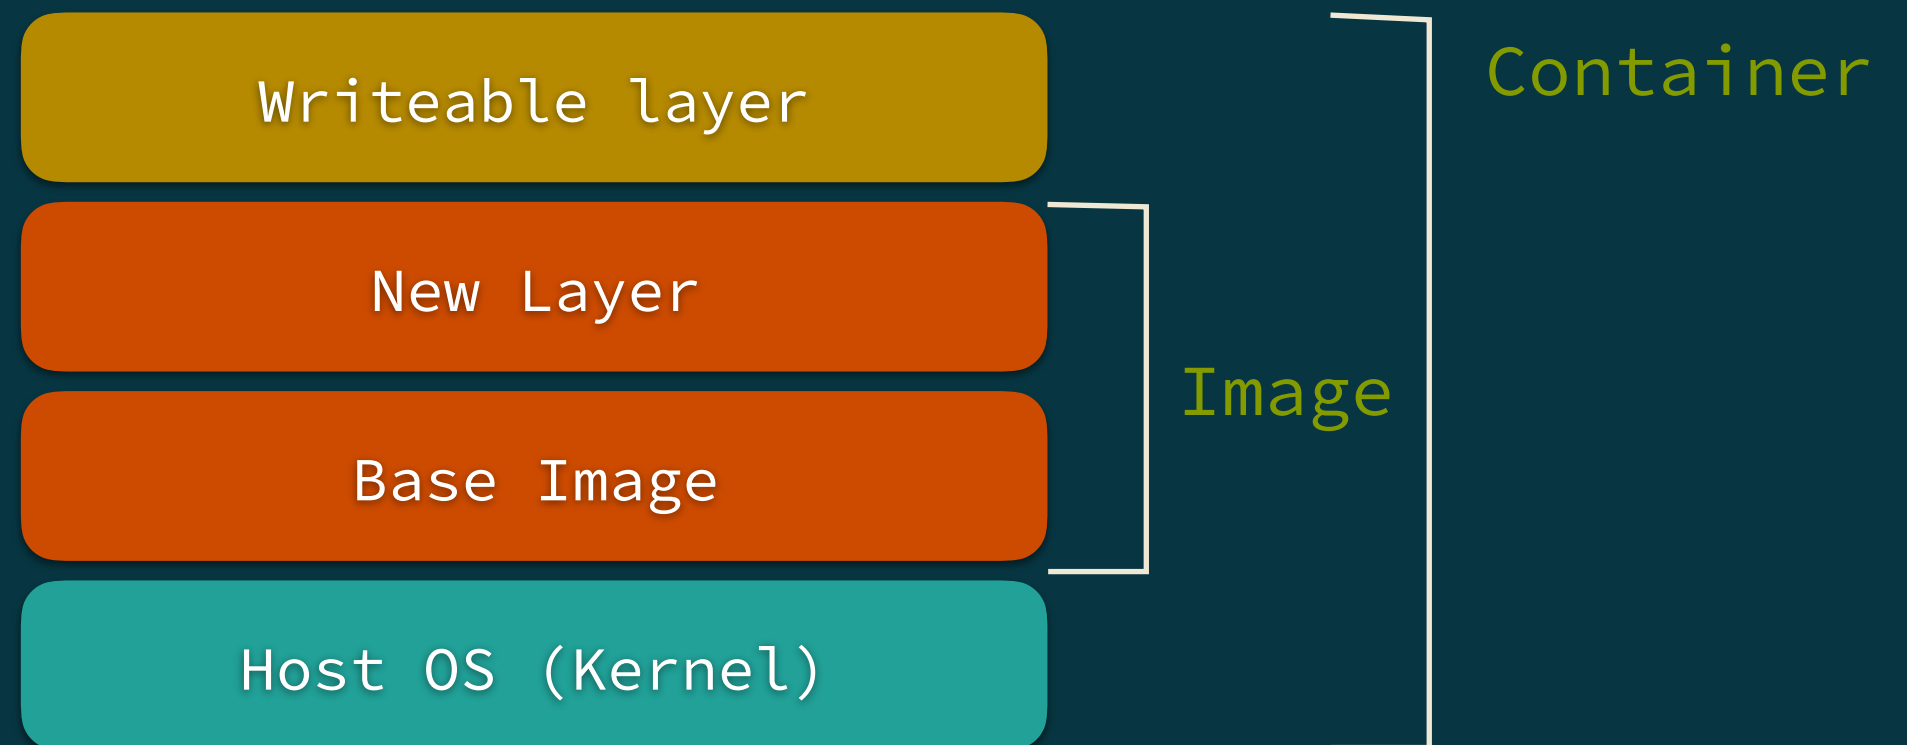
Image

Container







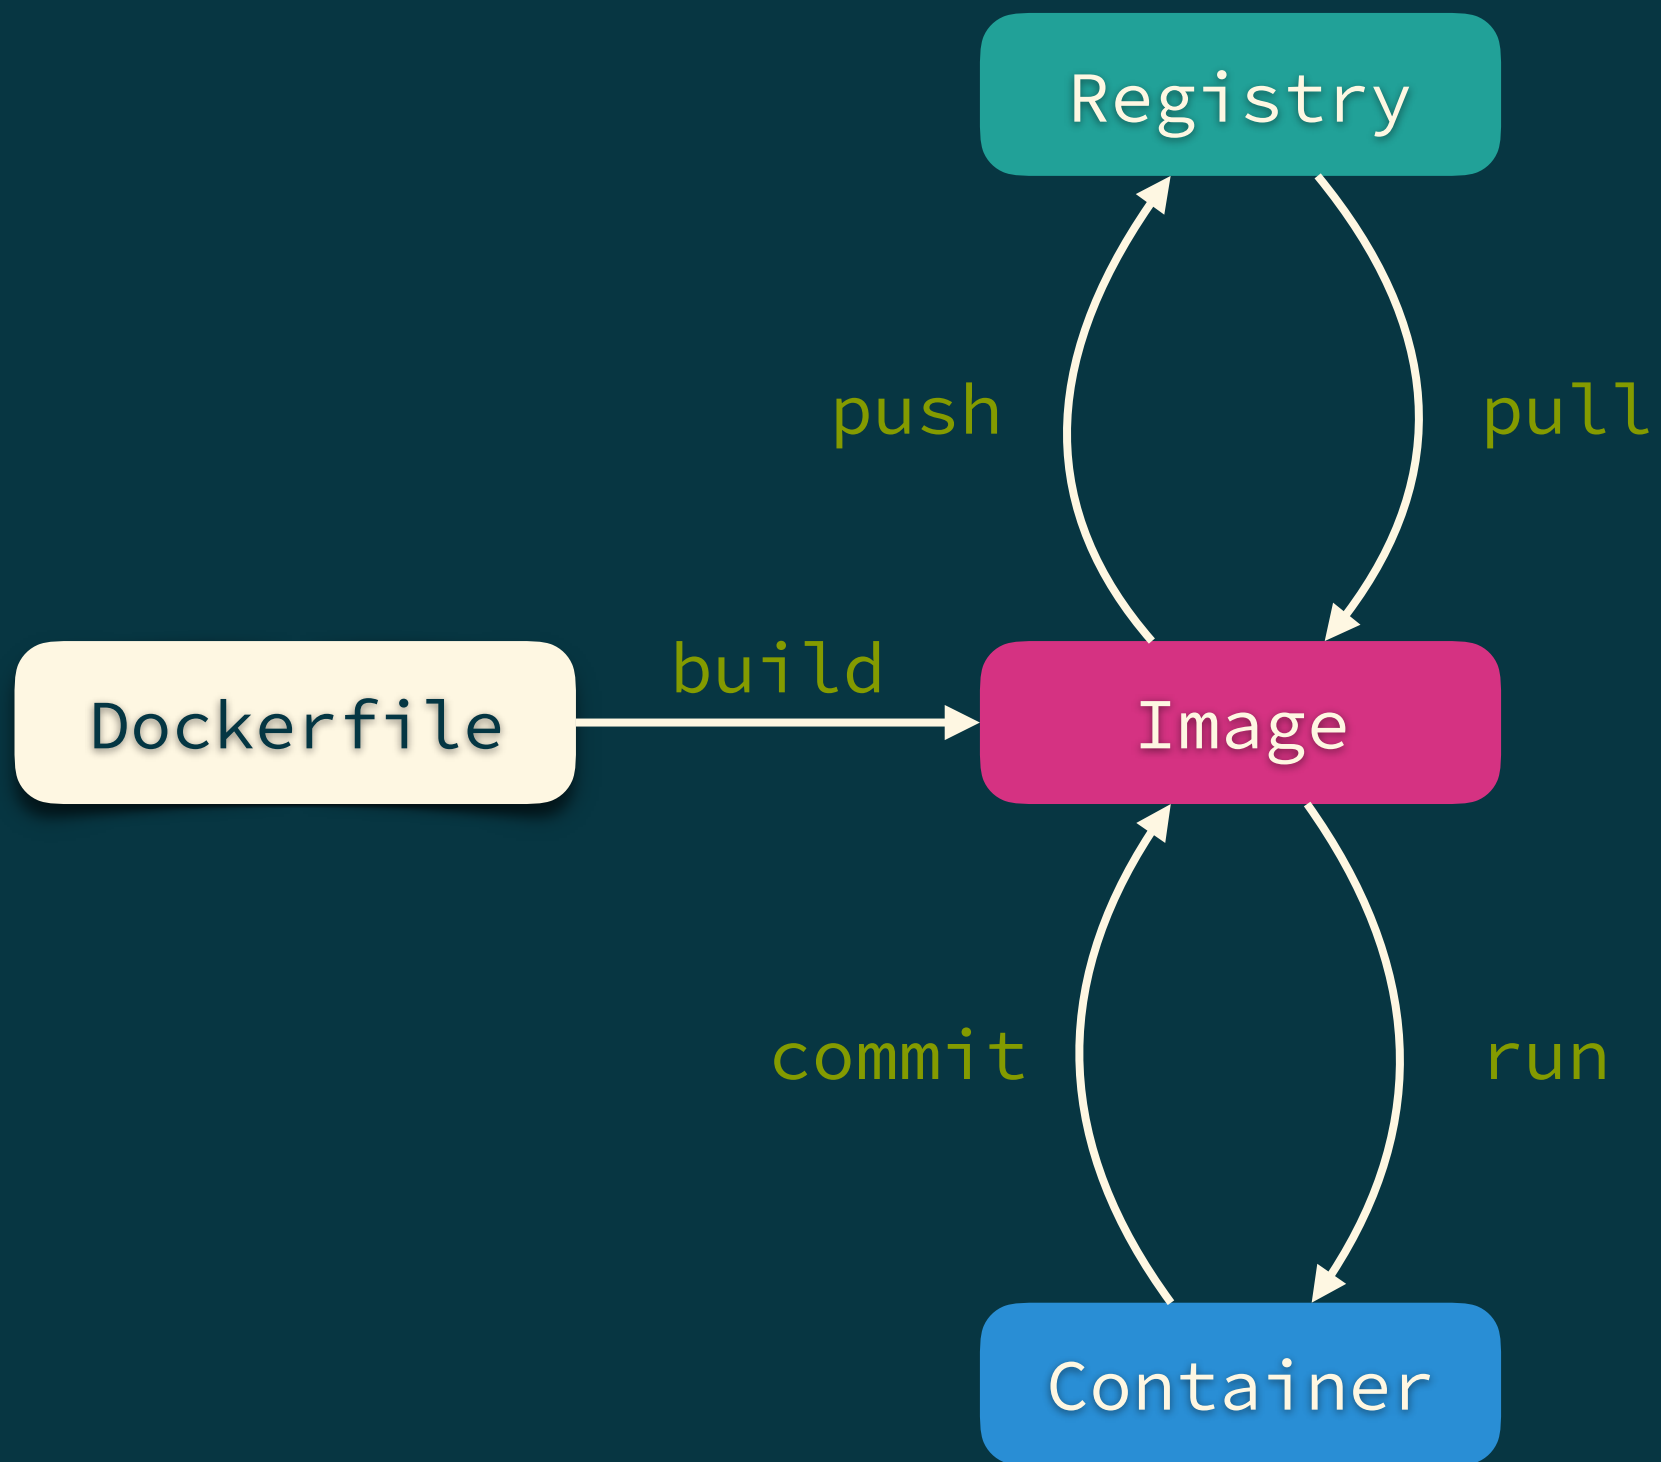
docker run











# DOCKERFILE

# DOCKERFILES

- A set of instructions to build a Docker image
- Plain text, version controlled
- Provides insight into the image needs/capabilities/intents

```
# sample Dockerfile
```


```
FROM openjdk:8u131-jre
```

```
RUN apt-get update \  
    && apt-get install -y netcat
```

```
COPY build/libs/app-fat.jar /var/app.jar
```

```
CMD ["java", "-jar", "/var/app.jar"]
```

# UNION FILE SYSTEM



**FROM** openjdk:8u131-jre

**RUN** apt-get update \  
 && apt-get install -y netcat

**COPY** build/libs/app-fat.jar /var/app.jar

**CMD** ["java", "-jar", "/var/app.jar"]







```
#!/usr/bin/env bash
```

```
last=alpine:3.8
```

```
for i in `seq 200`; do
```

```
    rm -f cid
```

```
    docker run --cidfile=cid $last touch file$i;
```

```
    docker commit `cat cid` tag$i;
```

```
    docker rm `cat cid`;
```

```
    last=tag$i;
```

```
done
```



```
#!/usr/bin/env bash
```

```
last=alpine:3.8
```

```
for i in `seq 200`; do
```

```
    rm -f cid
```

```
    docker run --cidfile=cid $last touch file$i;
```

```
    docker commit `cat cid` tag$i;
```

```
    docker rm `cat cid`;
```

```
    last=tag$i;
```

```
done
```

```
# Error response from daemon: max depth exceeded
```

```
# Unable to find image 'tag125:latest' locally
```



```
FROM alpine:3.8
```

```
RUN touch file1
```

```
RUN touch file2
```

```
# 125 times more
```

```
RUN touch file127
```



```
FROM alpine:3.8
```


```
RUN touch file1
```

```
RUN touch file2
```

```
# 125 times more
```

```
RUN touch file127
```

```
# Error response from daemon: max depth exceeded
```




**FROM** openjdk:8u131-jre

**RUN** apt-get update \  
&& apt-get install -y netcat

**COPY** build/libs/app-fat.jar /var/app.jar

**CMD** ...



FROM openjdk:8u131-jre

RUN apt-get update \  
&& apt-get install -y netcat

COPY build/libs/app-fat.jar /var/app.jar

CMD ...

**ONE CONCERN PER CONTAINER**

**@LOOSELYTYPED**



# (POLL - Single Choice)

## Base image usage

- We try and use as much as we can from public registries
- Outside of a handful of base images we build everything in-house
- Not concerned about this (yet)

**FROM**

# NOTES

- Implies “ancestry”
- **Has** to be the first line (Except if preceded by ARG)
- Has implications on WORKDIR, USER, ENTRYPOINT (and CMD), and ONBUILD, EXPOSE and other commands
  - Use "docker inspect" for this
- Create a base image with FROM scratch



- Pin down the exact tag (or even better the digest)
  - Do not use “latest” tag
- Choose your parent image wisely
  - Vet it!
  - Inspect ancestor images for USERS, PORTs, ENVs, VOLUMES, LABELs and anything that can be inherited
- Most likely you will build the lineage yourself internal to your team and organization

```
# Don't  
FROM alpine
```

```
# Do  
# Pin the version  
FROM alpine:3.8  
# OR Use ARG to set it at build time  
ARG version=3.8  
FROM alpine:${version}
```

**ARG**

# DO'S



- Use ARG for tweaking the build dynamically
  - Use them to set FROM, ENV, LABEL, RUN
- Default them appropriately

```
# Do
# Use default value
ARG AUTHOR="Raju Gandhi"
ARG BUILD_DATE
ARG VCS_REF
```

```
# default author
docker build --build-arg -t test .
# set at build time
docker build --build-arg AUTHOR="Solomon Hykes" -t test .
```



**ENV**

# DONT'S



- Put secrets or sensitive information in ENV variables
- Override parent image ENV's unless absolutely necessary

# DO'S



- Use them for documentation and modifying runtime behavior
  - They are baked in the final image
- Use `docker run <image-name> env`
  - Or `docker inspect`
- Default then appropriately
- Be cognizant of inherited ENV variables

```
ARG PROJECT_VERSION
```

```
# Do
```

```
# Default them if set dynamically
```

```
ENV PROJECT_VERSION ${PROJECT_VERSION:-2.3}
```

**LABEL**

# DONT'S



- Define individual labels separately

# DO'S



- Use them liberally
- Labels can see ARG variables. Use this!
  - BUILD\_NUMBER, GIT\_SHA
- Apply a standard convention
  - Build tooling on top of the conventions

```
ARG AUTHOR="Raju Gandhi"
```

```
ARG BUILD_DATE
```

```
ARG VCS_REF
```

```
# Don't
```

```
LABEL org.label-schema.author=$AUTHOR
```

```
LABEL org.label-schema.build-date=$BUILD_DATE
```

```
LABEL org.label-schema.vcs-ref=$VCS_REF
```

```
# Do
```

```
LABEL org.label-schema.author=$AUTHOR \  
    org.label-schema.build-date=$BUILD_DATE \  
    org.label-schema.vcs-ref=$VCS_REF
```



# **(POLL - Multiple Choice)**

## **Concerns with image size**

- Affects our build/runtimes
- Security concerns
- Disk usage for storage (local/registry/servers)
- Not a concern (yet)

**RUN**

# DON'TS



- Be cognizant of the effects (and drawbacks) of caching
- Do not do OS level upgrades (eg. `RUN dist-upgrade`)



- Group common operations
  - Clean up as well (reduces image sizes)
- Use multiline (\) to make PR / auditing easier

# Don't

**RUN** apt-get update

**RUN** apt-get install -y netcat

**RUN** apt-get clean

# Do

**RUN** apt-get update \  
    && apt-get install -y \  
    netcat \  
    && apt-get clean

**ENTRYPOINT/CMD**

# NOTES

- CMD can be overridden if the user supplies an argument to create or run
- You can (also) override ENTRYPOINT by explicitly supplying `--entrypoint` flag
- The default command run in a container is (ENTRYPOINT + CMD)

# DONT'S



- Avoid the “shell” form





- Use the “exec” form
  - Shell expansion will **not** happen!
- Use ENTRYPOINT and CMD together
- Use a “entrypoint-script”
  - Allows you to set error (-e) flags and traps
  - Your editor is syntax aware
  - Always “exec”

```
# Do
# Use ENTRYPOINT and CMD together
ENTRYPOINT [ "echo", "hello" ]
CMD [ "world" ]

# > docker build -t entrypoint-cmd .
# > docker run entrypoint-cmd
# hello world
# > docker run entrypoint-cmd raju
# hello raju
```

```
# Do
```

```
COPY entrypoint.sh /usr/local/bin/
```

```
ENTRYPOINT ["entrypoint.sh"]
```

```
CMD ["default"]
```

```
# entrypoint.sh
```

```
if [ "$1" = 'default' ]; then
```

```
    # do default thing here
```

```
    echo "Running default"
```

```
else
```

```
    echo "Running user supplied arg"
```

```
    # if the user supplied say /bin/bash
```

```
    exec "$@"
```

```
fi
```

**HEALTHCHECK**

# DONT'S



- Be too aggressive with `--interval` period
  - Especially if the check itself is expensive
- Use external tools (like `curl`) if you can

# DO'S



- Use a script that leverages the same runtime as your service
- For example, if you have a node or go service, write a health check using the same runtime
- Be cognizant of the overhead the health check introduces
- Experiment with combinations of interval/timeout/retries

# Avoid

HEALTHCHECK CMD curl --fail http://localhost:5000/ || exit 1

# Do

COPY healthcheck ./healthcheck

HEALTHCHECK --interval=1s \  
--timeout=1s \  
--start-period=2s \  
--retries=3 CMD [ "/healthcheck" ]

**ADD/COPY**



**.DOCKERIGNORE**

```
# example .dockerignore file
# ignore these folders
.git
build
!build/libs/*.jar # but NOT this
```

```
# ignore these files
.project
.gitignore
.dockerignore
```

```
# ignore all Docker files
Dockerfile*
docker-compose.yml
```

```
# ignore all markdown files (md) besides README.md
*.md
```

# DON'TS



- Avoid ADD
- Do not leave “residual” artifacts
- Be wary of using the array syntax
- Copy over all source in one fell swoop
  - Copy over source files separately and later on since they change often



- Instead of ADD
  - Combine COPY and RUN
  - OR RUN with wget/curl/tar/unzip
  - See DO'S under RUN
- Be mindful of what you put in the `.dockerignore` file



**USER**

# DON'TS



- Do not switch USER often
- Avoid using root

# DO'S



- Create a user (if you can) for your service
- Default the container to a non-root user if you can



```
FROM ubuntu:18.10
```

```
# Do
```

```
RUN groupadd -r app \  
    && useradd -r -g app appuser
```

```
USER appuser
```

**EXPOSE**

# DON'TS



- Avoid “docker run -P”

# DO'S



- Do document the ports your application needs exposed

# COMPILE VS RUNTIME



```
FROM alpine:3.8
```

```
RUN apk add --update \  
    tzdata \  
    && rm -rf /var/cache/apk/*
```

```
ARG TZ=America/Los_Angeles
```

```
RUN ln -snf \  
    /usr/share/zoneinfo/$TZ \  
    /etc/localtime && echo $TZ > /etc/timezone
```



FROM alpine:3.8

RUN apk add --update \  
tzdata \  
&& rm -rf /var/cache/apk/\*

ARG TZ=Americal/Los\_Angeles

ENV TZ \$TZ

COPY entrypoint.sh /usr/local/bin/

ENTRYPOINT ["entrypoint.sh"]

CMD ["default"]

```
#!/bin/sh
```

```
ln -snf \  
    /usr/share/zoneinfo/$TZ \  
    /etc/localtime && echo $TZ > /etc/timezone
```



- Use a combination of ARG and ENV with ENTRYPOINT/CMD
- Allows you to express what is at "compile" time versus "runtime"
- ARG/ENV are important parts of your image and containers API



**MULTI-STAGE BUILDS**

build

```
FROM openjdk:8u131-jdk as builder
WORKDIR /code
ADD . ./
RUN ["/gradlew", "shadowJar", "--no-daemon"]
```

run



cp



build

```
FROM openjdk:8u131-jre
RUN apt-get update \
    && apt-get install -y \
    netcat \
    && apt-get clean
COPY docker-workshop-0.0.1-SNAPSHOT-fat.jar \
    /var/app.jar
CMD ["java", "-jar", "/var/app.jar"]
```

# NOTES

- Allow you to do everything using docker containers
- Separates build environment from runtime, with the net effect of leaner production images
  - Keep secrets out of production images
  - Image builds are much faster because docker can leverage the cache

```
FROM openjdk:8u131-jdk as builder
WORKDIR /code
ADD . ./
RUN ["../gradlew", "shadowJar", "--no-daemon"]
```

```
FROM openjdk:8u131-jre
RUN apt-get update \
    && apt-get install -y \
    netcat \
    && apt-get clean
EXPOSE 8080
COPY --from=builder \
    /code/build/libs/docker-workshop-0.0.1-SNAPSHOT-fat.jar \
    /var/app.jar
CMD ["java", "-jar", "/var/app.jar"]
```

```
FROM openjdk:8u131-jdk as builder
WORKDIR /code
ADD . ./
RUN ["../gradlew", "shadowJar", "--no-daemon"]
```

```
FROM openjdk:8u131-jre
RUN apt-get update \
    && apt-get install -y \
    netcat \
    && apt-get clean
EXPOSE 8080
COPY --from=builder \
    /code/build/libs/docker-workshop-0.0.1-SNAPSHOT-fat.jar \
    /var/app.jar
CMD ["java", "-jar", "/var/app.jar"]
```

**THANKS!!**

# RESOURCES

- [Best practices for writing Dockerfiles](#)
- [Docker Registry V2](#)
- [Explaining Docker Image IDs](#)
- Dockerfiles for reference
  - [redis](#)
  - [Jenkins](#)
  - [Postgresql](#)